

PROJECT REPORT Milestone I

Linear regression :

For linear regression, we began our tests without giving any specific attention to the data (if not for the creation of a validation set which is done for every method). We tested the method for every λ between 0 and 100 and got a small variation in the results as it went from $MSE = 0.0510$ (for $\lambda = 0$) to $MSE = 0.0514$ (for $\lambda = 100$), with the minimum being slightly below 0.0510 (for $\lambda = 20$).

Realizing that we couldn't improve the MSE by modifying the λ , we decided to try adding a bias term to the data by adding an argument - bias_term in the main. The impact of the bias term was immediate as for the same λ 's than above, the MSE is divided by 10 (see Fig.1). We then decided to test more specific values of λ around $\lambda = 1$ as it was the best value on the interval $[0, 100]$. With the more precise test of λ (see Fig.2), we found that the best loss was achieved for $\lambda = 1.2$.

Best results ($\lambda = 1.2$):

Test loss = 0.0046

Fit time : 0.00018692016601562 seconds

Predict time : 0.0000030994415 seconds

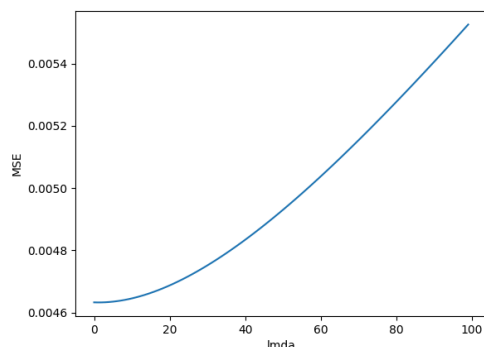


Fig.1 : MSE as a function of the lambda parameter ($\lambda \in [0, 100]$)

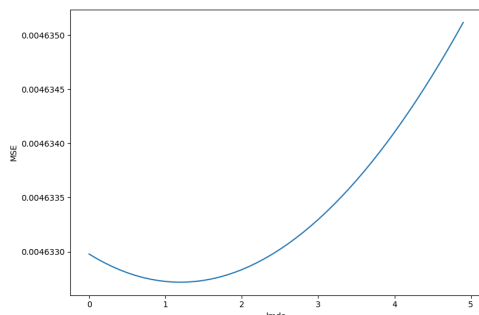


Fig.2 : MSE as a function of the lambda parameter ($\lambda \in [0, 5]$)

Logistic regression :

Similarly to the previous method, we started the test of logistic regression without processing the data. We tried for 5 different max_iters from 100 to 300 and for learning rates between 0.00001 and 0.1. This gave us 77.982% of accuracy at best, for $lr = 0.001$. We were stuck to this percentage, no matter what precision we had on the learning rates we tried. Hence, we decided again to add a bias term with the argument - bias_term . With the data biased, we finally got a better accuracy (see Fig.3). We could then conduct more precise tests around $lr = 0.001$, which led us to the conclusion that the best accuracy occurs when $\text{max_iters} = 100$ and $lr = 0.0045$ (see Fig.4).

Best results ($lr = 0.0045$, $\text{max_iters} = 100$):

Test set: accuracy = 87.156%

Fit time : 0.06351590156555176 seconds

Predict time : 0.0000572204589 seconds

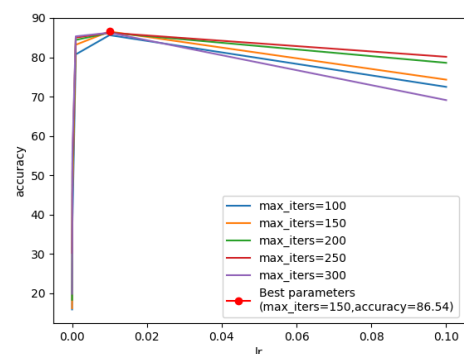


Fig.3 : accuracy as a function of the learning rate ($lr \in [0.00001, 0.1]$) for different max_iters

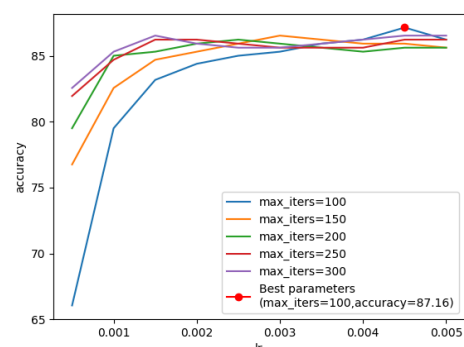


Fig.4 : accuracy as a function of the learning rate ($lr \in [0.0005, 0.005]$) for different max_iters

PROJECT REPORT Milestone I

KNN :

In this part, we calculate the k-nearest neighbors by first computing the Euclidean distance between a given example and all the training points. We opt for the Euclidean distance seen in class, due to its simplicity and effectiveness, but alternatives like the chi-square distance could have also been utilized. After computing distances, we identify the k nearest elements. K serves as a hyperparameter, allowing us to specify the number of neighbors to consider. Increasing the value of K can result in more generalized patterns. However, it's important to note that excessively large values of K may lead to over-smoothing and potentially poorer performance on unseen data.

That is why when first testing values of K , for KNN as a Regression method, we decided to choose the interval $[1,10]$.

Surprisingly, it was a decreasing function and the best K parameter returned was **10**. We then decided to extend the interval to $[1,20]$ and then $[1,30]$ but it always seemed like as K grows, the loss was decreasing. Finally, we decided to test on a very large interval to be fixed about our hypothesis and we choose to test K on the interval $[1,100]$. The results are clear, the loss is indeed decreasing as K grows but converges to **0.0046** for $K \geq 30$ (see Fig.5).

For KNN as a Classification method, we again started by testing for a small interval of K (see Fig.6). For this task, the results fitted our expectations as for $K=1$, the method was clearly overfitting (Train accuracy almost equal to **100%** but just **83%** for test data) and the best K parameter ($K=6$) gave us almost equivalent accuracy on both training and test data (test accuracy was **88.685%**).

KNN Regression best results:

Test loss = 0.0055

Fit time : 0.45464015007019043 seconds

Predict time : 0.052410125732421875 seconds

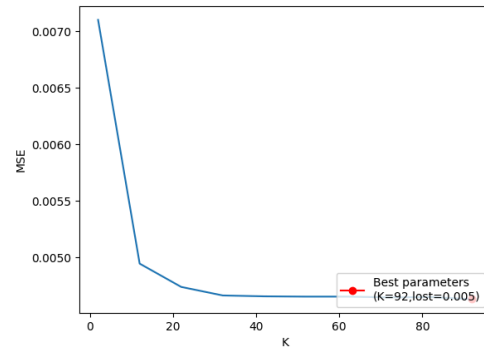


Fig.5 : MSE as a function of the K parameter ($K \in [1,100]$)

KNN Classification best results:

Test set: accuracy = 86.239%

Fit time : 0.4482560157775879 seconds

Predict time : 0.05003499984741211 seconds

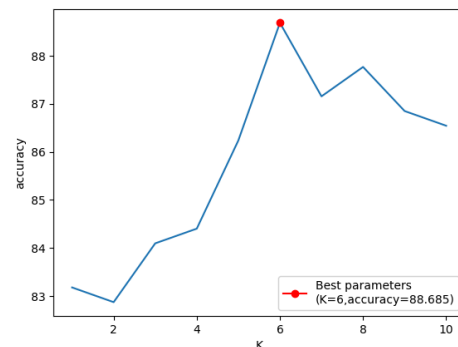


Fig.6 : accuracy as a function of K parameter ($K \in [1,10]$)